

INTRODUCTION TO SUPERCOMPUTING

TMA4280 · Project I

This project is mandatory, counts for 10% of the final grade and can be done in pairs or alone.

Instructions

The deadline is set on the **1st March 2017**.

The deliverable consists of:

1. a report describing your solution, handed out by email in PDF format,
2. a GIT repository containing the source code developed to perform the computations.

Practical requirements regarding the code:

1. it should be hosted on Github,
2. it should be structured in different subdirectories addressing the different questions and results gathered in text files,
3. it must be written in C/C++ (without use of `std::vector`) or FORTRAN,
4. it must use double precision,
5. it must compile and run using Makefile targets,
6. results presented in the report have to be reproducible.

As soon as you have decided whether you want to work in pairs or alone:

- create a Github repository (one per pair) named TMA4280LABS,
- send your name(s) and the link to the repository by email.

Description

The project deals with the computation of π using different approximations and their parallelization using MPI and OpenMP.

Considered approximations are based on the development in series,

$$S_n = \sum_{i=1}^n v_i.$$

with v_i depending on the chosen method.

1. Riemann zeta function $\zeta(s)$ with $s = 2$: Each term is defined as:

$$v_i = \frac{1}{i^2}, \quad i = 1, \dots, n$$

and

$$S = \lim_{n \rightarrow \infty} S_n = \frac{\pi^2}{6}.$$

2. Machin Formula: Given $x \in [-1, 1]$,

$$v_i(x) = (-1)^{i-1} \frac{x^{2i-1}}{2i-1}$$

and

$$S(x) = \lim_{n \rightarrow \infty} S_n(x) = \arctan(x)$$

with

$$\frac{\pi}{4} = 4 \arctan\left(\frac{1}{5}\right) - \arctan\left(\frac{1}{239}\right)$$

In all the questions, the name in bold and enclosed in parenthesis denotes the subdirectory where the implementation should be located. You may factor common code in a library if you want.

Question 1. Serial implementation. Write a serial program implementing the computation of π for a given n read from the command line for:

- Method 1. add a program in (**zeta0**)
- Method 2. add a program in (**mach0**)

Question 2. Unit test. Every development should come with unit testing to check the logic of the implementation. Such tests should execute quickly and compare a computed value against an expected value. Implement a small test comparing the value of each series with $n = 3$.

- Method 1. add a unit test in (**zeta0**)
- Method 2. add a unit test in (**mach0**)

The test should be implemented in a simple function (no unit test framework required) and executed with `make utest`. Why do you think such test may be useful when parallelizing a computational code?

Question 3. Verification test. When the logic of the implementation is tested then the mathematical properties of the algorithm should be assessed: convergence, stability. Such tests should execute fairly quickly and verify the behaviour of the algorithm as compared to the theory if possible. Implement a small test computing the error $|\pi - \pi_n|$ for $n = 2^k$ with $k = 1, \dots, 24$:

- Method 1. add a verification test in (**zeta0**)
- Method 2. add a verification test in (**mach0**)

The test should be executed with `make vtest` and results should be saved in a file. Comment on the obtained results.

Question 4. Data distribution. As we are interested in computing the sum of all vector elements v_i numerically, we will work under the constraint that the values should be put in a vector before being summed. The suggest program deliberately relies on partitioning and distribution of the data by process zero.

Process zero **only** should be responsible for:

- generating the vector elements,
- partitioning the vector in a way that the problem is load-balanced,
- and distribute the elements to all the processes.

Each process will work on the received data.

- Method 1. add a program in (**zeta1**)
- Method 2. add a program in (**mach1**)

Can you comment on the limitation of such approach for the data distribution and a possible improvement?

Question 5. MPI implementation. Modify further the program to compute the approximation of π using both methods such that each process:

- computes a partial sum from its data,

- then all the partial sums should be added together on the root process,
- and then global sum is printed on the standard output by the root process.

Only the root process holds the final value. Report the error $|\pi - \pi_n|$ in double precision for different values of n and the wall time, for different number of MPI processes which are powers of two. The program should contain an assertion and fail if the number of processes is not a power of two, as a design constraint.

- Method 1. modify the program and add a test in (**zeta1**)
- Method 2. modify the program and add a test in (**mach1**)

Plot the error and the timings. Which MPI calls were convenient and/or necessary to use? Can you comment on the methodology used for computing the wall time?

Question 6. Analysis

- Compare the errors from the single-process program and the multi-process program for $P = 2$ and $P = 8$. Should the answer be the same in all cases? Exactly, or approximately? Can you explain why?

Question 7. Global reduction. Modify the final step, the reduction with `MPI_SUM`, such that all processes store the global sum: first by using an MPI function and then by implementing the recursive-doubling sum.

- Add a program in (**reduc**)

Do a small scaling study, what do you observe?

Question 8. OpenMP implementation. Make the necessary changes needed to use shared memory parallelization with OpenMP.

- Method 1. add a program in (**zeta2**)
- Method 2. add a program in (**mach2**)

Perform the same analysis as for the *MPI implementation*.

Question 9. Hybrid MPI/OpenMP implementation. Confirm that your program also works when using OpenMP and MPI in combination.

Question 10. Discussion.

Compare the memory requirement per process for the single-process program and the multi-process program when $n \gg 1$.

How many floating point operations are needed to generate the vector v ?

How many are needed to compute S_n ?

Is the multi-process program load balanced?

Question 11. Conclusion. Do you consider parallel processing attractive for solving this problem? Explain why.